



DEPARTMENT OF PHILOSOPHY,
LINGUISTICS AND THEORY OF SCIENCE

ON THE GRAMMAR OF PROOF

Warrick Macmillan

Master's Thesis:	30 credits
Programme:	Master's Programme in Language Technology
Level:	Advanced level
Semester and year:	Fall, 2021
Supervisor	Aarne Ranta
Examiner	(name of the examiner)
Report number	(number will be provided by the administrators)
Keywords	Grammatical Framework, Natural Language Generation,

Abstract

Brief summary of research question, background, method, results...

Preface

Acknowledgements, etc.

Contents

1335

Add section numbers to sections

Introduction

The central concern of this thesis is the syntax of mathematics, programming languages, and their respective mutual influence, as conceived and practiced by mathematicians and computer scientists. From one vantage point, the role of syntax in mathematics may be regarded as a 2nd order concern, a topic for discussion during a Fika, an artifact of ad hoc development by the working mathematician whose real goals are producing genuine mathematical knowledge. For the programmers and computer scientists, syntax may be regarded as a matter of taste, with friendly debates recurring regarding the use of semicolons, brackets, and white space. Yet, when viewed through the lens of the propositions-as-types paradigm, these discussions intersect in new and interesting ways. When one introduces a third paradigm through which to analyze the use of syntax in mathematics and programming, namely linguistics, I propose what some may regard as superficial detail, indeed becomes a central paradigm raising many interesting and important questions.

Beyond Computational Trinitarianism

The doctrine of computational trinitarianism holds that computation manifests itself in three forms: proofs of propositions, programs of a type, and mappings between structures. These three aspects give rise to three sects of worship: Logic, which gives primacy to proofs and propositions; Languages, which gives primacy to programs and types; Categories, which gives primacy to mappings and structures.?

We begin this discussion of the three relationships between three respective fields, mathematics, computer science, and logic. The aptly named trinity, shown in Figure 4, are related via both *formal* and *informal* methods. The propositions as types paradigm, for example, is a heuristic. Yet it also offers many examples of successful ideas translating between the domains. Alternatively, the interpretation of a Type Theory(TT) into a category theory is incredibly *formal*.

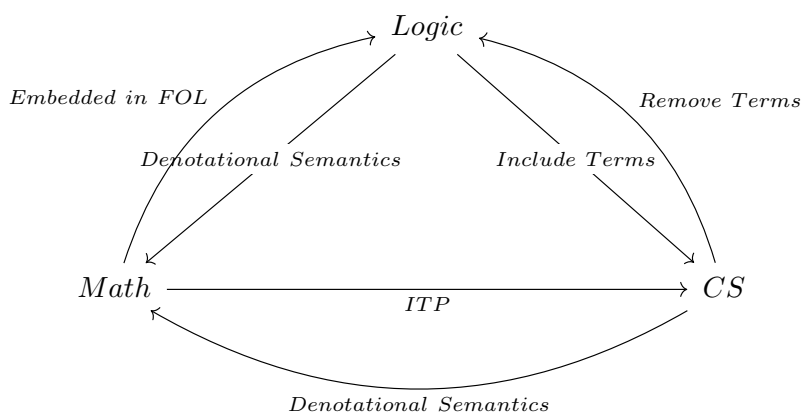


Figure 1: The Holy Trinity

We hope this thesis will help clarify another possible dimension in this diagram, that of Linguistics, and call it the “holy tetrahedron”. The different vertices also resemble religions in their own right, with communities convinced that they have a canonical perspective on foundations and the essence of mathematics. Questioning the holy trinity is an act of a heresy, and it is the goal of this thesis to be a bit heretical by including a much less well understood perspective which provides additional challenges and insights into the trinity.

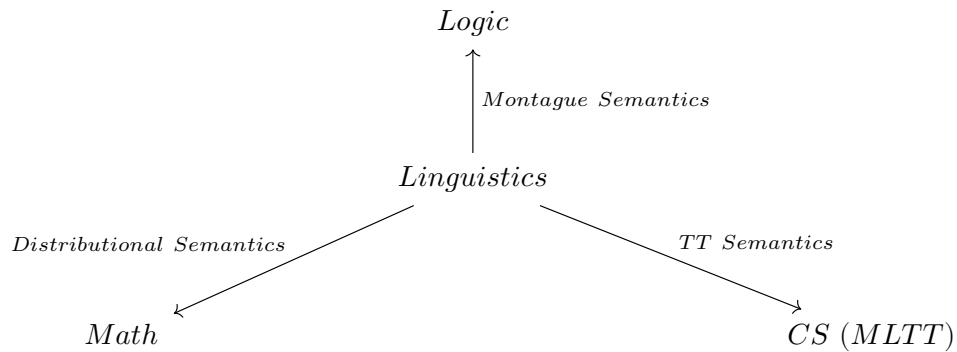


Figure 2: Formal Semantics

One may see how the trinity give rise to *formal* semantic interpretations of natural language in Figure 5. Semantics is just one possible linguistic phenomenon worth investigating in these domains, and could be replaced by other linguistic paradigms. This thesis is alternatively concerned with syntax.

Finally, as in Figure 3, we can ask : how does the trinity embed into natural language? These are the most *informal* arrows of tetrahedron, or at least one reading of it. One can analyze mathematics using linguistic methods, or try to give a natural language justification of Intuitionistic Type Theory (ITT) using Martin-Löf’s meaning explanations.

In this work, we will see that there are multiple GF grammars which model some subset of each member of the trinity. Constructing these grammars, and asking how they can be used in applications for mathematicians, logicians, and computer scientists is an important practical and philosophical question. Therefore we hope this attempt at giving the language of mathematics, in particular how propositions and proofs are expressed and thought about in that language, a stronger foundation.

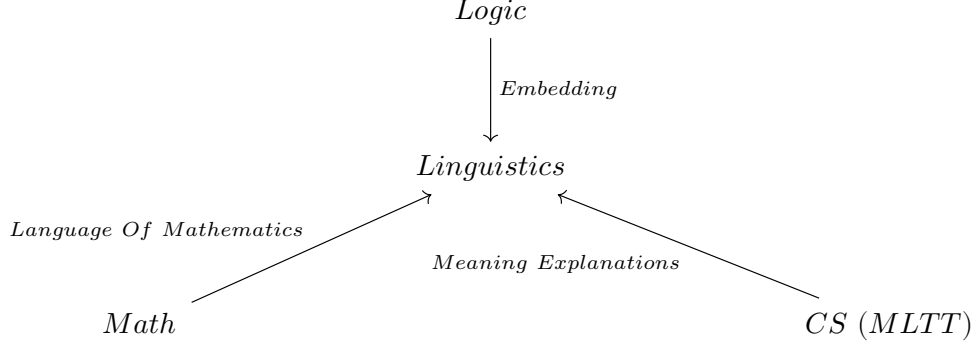


Figure 3: Interpretations of Natural Language

What is a Homomorphism?

To get a feel for the syntactic paradigm we explore in this thesis, let us look at a basic mathematical example: that of a group homomorphism as expressed in by a variety of somewhat randomly sampled authors.

Definition 1 *In mathematics, given two groups, $(G, *)$ and (H, \cdot) , a group homomorphism from $(G, *)$ to (H, \cdot) is a function $h : G \rightarrow H$ such that for all u and v in G it holds that*

$$h(u * v) = h(u) \cdot h(v)$$

Definition 2 *Let $G = (G, \cdot)$ and $G' = (G', *)$ be groups, and let $\phi : G \rightarrow G'$ be a map between them. We call ϕ a **homomorphism** if for every pair of elements $g, h \in G$, we have*

$$\phi(g * h) = \phi(g) \cdot \phi(h)$$

Definition 3 *Let G, H , be groups. A map $\phi : G \rightarrow H$ is called a group homomorphism if*

$$\phi(xy) = \phi(x)\phi(y) \text{ for all } x, y \in G$$

(Note that xy on the left is formed using the group operation in G , whilst the product $\phi(x)\phi(y)$ is formed using the group operation H .)

Definition 4 *Classically, a group is a monoid in which every element has an inverse (necessarily unique).*

We inquire the reader to pay attention to nuance and difference in presentation that is normally ignored or taken for granted by the fluent mathematician, ask which definitions feel better, and how the reader herself might present the definition differently.

If one want to distill the meaning of each of these presentations, there is a significant amount of subliminal interpretation happening very much analogous to our innate linguistic usage. The inverse and identity are discarded, even though they are necessary data when defining a group. The order of presentation of information is inconsistent, as well as the choice to use symbolic or natural language information. In Definition 3, the group operation is used implicitly, and its clarification a side remark.

Details aside, these all mean the same thing - don't they? This thesis seeks to provide an abstract framework to determine whether two linguistically nuanced presentations mean the same thing via their syntactic transformations. Obviously these meanings are not resolvable in any kind of absolute sense, but at least from a translational sense. These syntactic transformations come in two flavors : parsing and linearization, and are natively handled by a Logical Framework (LF) for specifying grammars : Grammatical Framework (GF).

We now show yet another definition of a group homomorphism formalized in the Agda programming language:

```
isGroupHom : (G : Group {ℓ}) (H : Group {ℓ'}) (f : ( G ) → ( H )) → Type _
isGroupHom G H f = (x y : ( G )) → f (x G.+ y) ≡ (f x H.+ f y) where
  module G = GroupStr (snd G)
  module H = GroupStr (snd H)

record GroupHom (G : Group {ℓ}) (H : Group {ℓ'}) : Type (ℓ-max ℓ ℓ') where
  constructor grouphom

  field
    fun : ( G ) → ( H )
    isHom : isGroupHom G H fun
```

This actually *was* the Cubical Agda implementation of a group homomorphism sometime around the end of 2020. We see that, while a mathematician might be able to infer the meaning of some of the syntax, the use of levels, the distinguishing between `isGroupHom` and `GroupHom` itself, and many other details might obscure what's going on.

We provide the current, as of May 2021, definition via Cubical Agda. One may witness a significant number of differences from the previous version : concrete syntax differences via changes in camel case, new uses of `Group` vs `GroupStr`, as well as, most significantly, the identity and inverse preservation data not appearing as corollaries, but part of the definition. Additionally, we had to refactor the

commented lines to those shown below to be compatible with our outdated version of cubical. These changes would not just be interesting to look at from the author of the libraries's perspective, but also syntactically.

```

record IsGroupHom {A : Type ℓ} {B : Type ℓ'}
  (M : GroupStr A) (f : A → B) (N : GroupStr B)
  : Type (ℓ-max ℓ ℓ')
  where

    -- Shorter qualified names
  private
    module M = GroupStr M
    module N = GroupStr N

  field
    pres· : (x y : A) → f (M._+_ x y) ≡ (N._+_ (f x) (f y))
    pres1 : f M.0g ≡ N.0g
    presinv : (x : A) → f (M._ x) ≡ N._ (f x)
    -- pres· : (x y : A) → f (x M.· y) ≡ f x N.· f y
    -- pres1 : f M.1g ≡ N.1g
    -- presinv : (x : A) → f (M.inv x) ≡ N.inv (f x)

  GroupHom' : (G : Group {ℓ}) (H : Group {ℓ'}) → Type (ℓ-max ℓ ℓ')
  -- GroupHom' : (G : Group ℓ) (H : Group ℓ') → Type (ℓ-max ℓ ℓ')
  GroupHom' G H = Σ[ f ∈ (G .fst → H .fst) ] IsGroupHom (G .snd) f (H .snd)

```

While the last two definitions may carry degree of comprehension to a programmer or mathematician not exposed to Agda, it is certainly comprehensible to a computer : that is, it typechecks on a computer where Cubical Agda is installed. While GF is designed for multilingual syntactic transformations and is targeted for natural language translation, it's underlying theory is largely based on ideas from the compiler communities. A cousin of the BNF Converter (BNFC), GF is fully capable of parsing programming languages like Agda! And while the Agda definitions are just another concrete syntactic presentation of a group homomorphism, they are distinct from the natural language presentations above in that the colors indicate it has indeed type checked.

While this example may not exemplify the power of Agda's type-checker, it is of considerable interest to many. The type-checker has merely assured us that `GroupHom(')` are well-formed types - not that we have a canonical representation of a group homomorphism. The type-checker is much more useful than is immediately evident: it delegates the work of verifying that a proof is correct, that is, the work of judging whether a term has a type, to the computer. While it's of practical concern is immediate to any exploited grad student grading papers late on a Sunday night, its theoretical concern has led to many recent developments in modern mathematics. Thomas Hales solution to the Kepler Conjecture was seen as unverifiable by those reviewing it, and this led to Hales outsourcing the verification to

Interactive Theorem Provers (ITPs) HOL Light and Isabelle. This computer delegated verification phase led to many minor corrections in the original proof which were never spotted due to human oversight.

Fields medalist Vladimir Voevodsky had the experience of being told one day his proof of the Milnor conjecture was fatally flawed. Although the leak in the proof was patched, this experience of temporarily believing much of his life's work invalidated led him to investigate proof assistants as a tool for future thought. Indeed, this proof verification error was a key event that led to the Univalent Foundations Project ?.

While Agda and other programming languages are capable of encoding definitions, theorems, and proofs, they have so far seen little adoption, and in some cases treated with suspicion and scorn by many mathematicians. This isn't entirely unfounded : it's a lot of work to learn how to use Agda or Coq, software updates may cause proofs to break, and the inevitable imperfections we humans are prone to instilled in these tools . Besides, Martin-Löf Type Theory, the constructive foundational project which underlies these proof assistants, is often misunderstood by those who dogmatically accept the law of the excluded middle as the word of God.

It should be noted, the constructivist rejects neither the law of the excluded middle, nor ZFC. She merely observes them, and admits their handiness in certain citations. Excluded middle is indeed a helpful tool as many mathematicians may attest. The contention is that it should be avoided whenever possible - proofs which don't rely on it, or it's corollary of proof by contradiction, are much more amenable to formalization in systems with decidable type checking. And ZFC, while serving the mathematicians of the early 20th century, is lacking when it comes to the higher dimensional structure of n-categories and infinity groupoids.

What these theorem provers give the mathematician is confidence that her work is correct, and even more importantly, that the work which she takes for granted and references in her work is also correct. The task before us is then one of religious conversion. And one doesn't undertake a conversion by simply by preaching. Foundational details aside, this thesis is meant to provide a blueprint for the syntactic reformation that must take place.

We don't insist a mathematician relinquish the beautiful language she has come to love in expressing her ideas. Rather, it asks her to make a hypothetical compromise for the time being, and use a Controlled Natural Language (CNL) to develop her work. In exchange she'll get the confidence that Agda provides. Not only that, she'll be able to search through a library, to see who else has possibly already postulated and proved her conjecture. A version of this grandiose vision is explored in The Formal Abstracts Project ?, and it should practically motivate work.

Practicalities aside, this work also attempts to offer a nuanced philosophical perspective on the matter by exploring why translation of mathematical language, despite it's seemingly structured form, is difficult. We note that the natural language definitions of monoid differ in form, but also in pragmatic content. How one expresses formalities in natural language is incredibly diverse, and Definition 4

as compared with the prior homomorphism definitions is particularly poignant in demonstrating this. These differ very much in nature to the Agda definitions - especially pragmatically. The differences between the Cubical Agda definitions may be loosely called pragmatic, in the sense that the choice of definitions may have downstream effects on readability, maintainability, modularity, and other considerations when trying to write good code, in a burgeoning area known as proof engineering.

A pragmatic treatment of the language of mathematics is the golden egg if one wishes to articulate the nuance in how the notions proposition, proof, and judgment are understood by humans. Nonetheless, this problem is just now seeing attention. We hope that the treatment of syntax in this thesis, while a long ways away from giving a pragmatic account of mathematics, will help pave the way there.

Perspectives

...when it comes to understanding the power of mathematical language to guide our thought and help us reason well, formal mathematical languages like the ones used by interactive proof assistants provide informative models of informal mathematical language. The formal languages underlying foundational frameworks such as set theory and type theory were designed to provide an account of the correct rules of mathematical reasoning, and, as Gödel observed, they do a remarkably good job. But correctness isn't everything: we want our mathematical languages to enable us to reason efficiently and effectively as well. To that end, we need not just accounts as to what makes a mathematical argument correct, but also accounts of the structural features of our theorizing that help us manage mathematical complexity.?

Linguistic and Programming Language Abstractions

The key development of this thesis is to explore the formal and informal distinction of presenting mathematics as understood by mathematicians and computer scientists by means of rule-based, syntax oriented machine translation.

Computational linguistics, particularly those in the tradition of type theoretical semantics?, gives one a way of comparing natural and programming languages. Type theoretical semantics is concerned with the semantics of natural language in the logical tradition of Montague, who synthesized work in the shadows of Chomsky ? and Frege ?. This work ended up inspiring the GF system, a side effect of which was to realize that machine translation was possible as a side effect of this abstracted view of natural language semantics. Indeed, one such description of GF is that it is a compiler tool applied to domain specific machine translation. We may compare the “compiler view” of PLs and the “linguistics view” of NLs, and interpolate this comparison to other general phenomenon in the respective domains.

We will reference these programming language and linguistic abstraction ladders, and after viewing Figure 4, the reader should examine this comparison with her own knowledge and expertise in mind. These respective ladders are perhaps the most important lens one should keep in mind while reading this thesis. Importantly, we should observe that the PL dimension, the left diagram, represents synthetic processes, those which we design, make decisions about, and describe formally. Alternatively, the NL abstractions on the right represent analytic observations. They are therefore subject to different, in some ways orthogonal, constraints.

The linguistic abstractions are subject to empirical observations and constraints, and this diagram only serves as an atlas for the different abstractions and relations between these abstractions, which may be subject to modifications depending on the linguist or philosopher investigating such matters. The PL abstractions as represented, while also an approximation, serves as an actual high altitude blueprint for the design of programming languages. While the devil is in the de-

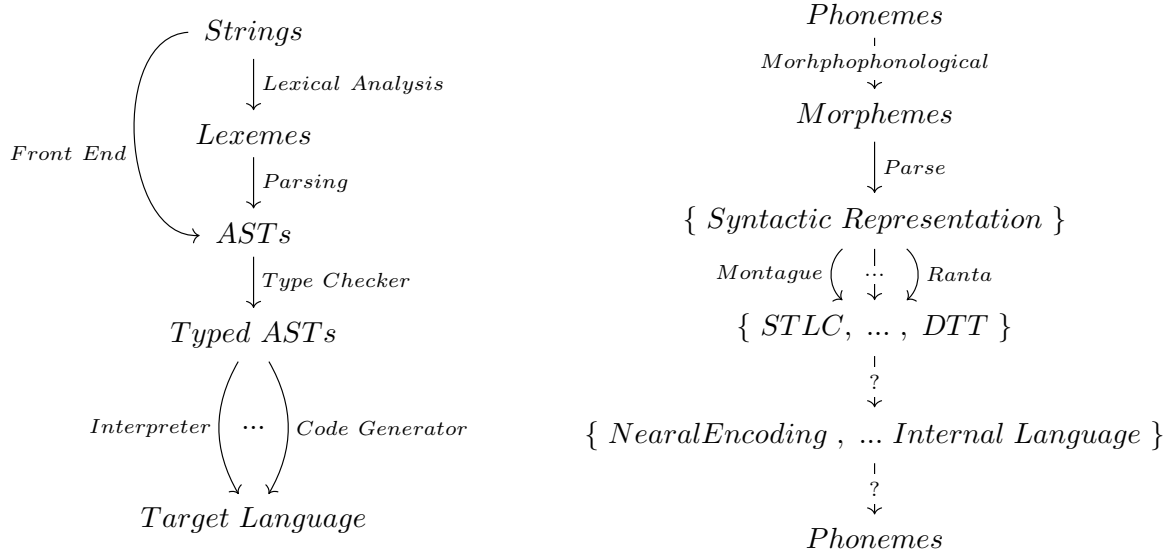


Figure 4: PL (left) and NL (right) Abstraction Ladders

tails and this view is greatly simplified, the representation of PL design is unlikely to create angst in the computer science communities. The linguistic abstractions are at the intersection of many fascinating debates between linguists, and there is certainly nothing close to any type of consensus among linguists which linguistic abstractions, as well as their hierarchical arrangement, are more practically useful, theoretically compelling, or empirically testable.

There are also many relevant concerns not addressed in either abstraction chain that are necessary to give a more comprehensive snapshot. For instance, we may consider intrinsic and extrinsic abstractions that diverge from the idealized picture. In PL extrinsic domain, we can inquire about

- systems with multiple interactive programming language
- how the programming languages behave with respect to given programs
- embedding programming languages into one another

Alternatively, intrinsic to a given PL, there picture is also not so clear. Agda, for example, requires the evaluation of terms during typechecking. It is implemented with 4.5 different stages between the syntax written by the programmers and the “fully reflected Abstract Syntax Tree (AST)” ?. But this example is perhaps an outlier, because Agda’s type-checker is so powerful that the design, implementation, and use of Agda revolves around it, (which, ironically, is already called during the parsing phase). It is not anticipated that floating point computation, for instance, would ever be considered when implementing new features of Agda, at least not for the foreseeable future. Indeed, the ways Agda represents ASTs were an obstacle encountered doing this work, because deciding which parsing stage one should connect to the Portable Grammar Format (PGF) embedding is nontrivial.

Let’s zoom in a little and observe the so-called front-end part of the compiler. Displayed in Figure 5 is the highest possible overview of GF. This is a deceptively




Figure 5: GF in a nutshell

simple depiction of such a powerful and intricate system. What makes GF so compelling is its ability to translate between inductively defined languages that type theorists specify and relatively expressive fragments of natural languages, via the composition of GF’s parsing and linearization capabilities. It is in some sense the attempt to overlay the abstraction ladders at the syntactic level and semantic led to this development.

For natural language, some intrinsic properties might take place, if one chooses, at the neurological level, where one somehow can contrast the internal language (i-language) with the mechanism of externalization (generally speech) as proposed by Chomsky ?. Extrinsic to the linguistic abstractions depicted, pragmatics is absent. . The point is to recognize their are stark differences between natural languages and programming languages which are even more apparent when one gets to certain abstractions. Classifying both programming languages as languages is best read as an incomplete (and even sometimes contradictory) metaphor, due to perceived similarities (of which there are ample).

Nonetheless, the point of this thesis is to take a crack at that exact question : how can one compare programming and natural languages, in the sense that a natural language, when restricted to a small enough (and presumably well-behaved) domain, behaves as a programming language. Simultaneously, we probe the topic of Natural Language Generation (NLG). Given a logic or type system with some theory inside (say arithmetic over the naturals), how do we not just find a natural language representation which interprets our expressions, but also does so in a way that is linguistically coherent in a sense that a competent speaker can make sense of it in a facile way.

The specific linguistic domain we focus on, that of mathematics, is a particular sweet spot at the intersection of these natural and formal language spaces. It should be noted that this problem, that of translating between *formal* and *informal* mathematics as stated, is both vague and difficult. It is difficult in both the practical sense, that it may be either of infeasible complexity or even perhaps undecidable, but it is also difficult in the philosophical sense. One may entertain the prospect of syntactically translated mathematics may a priori may deflate its effectiveness or meaningfulness. Like all collective human endeavors, mathemat-

ics is a historical construction - that is, its conventions, notations, understanding, methodologies, and means of publication and distribution have all been in a constant flux. There is no consensus on what mathematics is, how it is to be done, and most relevant for this treatise, how it is to be expressed.

Historically, mathematics has been filtered of natural language artifacts, culminating in some sense with Frege's development of a formal proof. A mathematician often never sees a formal proof as it is treated in Logic and Type Theory. We hope this work helps with a new foundational mentality, whereby we try to bring natural language back into mathematics in a controlled way, or at least to bridge the gap between our technologies, specifically injecting ITPs into a mathematicians toolbox.

We present a sketch of the difference of this so-called formal/informal distinction. Mathematics, that is mathematical constructions like numbers and geometrical figures, arose out of ad-hoc needs as humans cultures grew and evolved over the millennia. Indeed, just like many of the most interesting human developments of which there is a sparsely documented record until relatively recently, it is likely to remain a mystery what the long historical arc of mathematics could have looked like in the context of human evolution. And while mathematical intuitions precede mathematical constructions (the spherical planet precedes the human use of a ruler compass construction to generate a circle), we should take it as a starting point that mathematics arises naturally out of our linguistic capacity. This may very well not be the case, or at least not universally so, but it is impossible to imagine humans developing mathematical constructions elaborating anything particularly general without linguistic faculties. Despite whatever empirical or philosophical dispute one takes with this linguistic view of mathematical abilities, we seek to make a first order approximation of our linguistic view for the sake of this work. The discussion around mathematics relation to linguistics generally, regardless of the stance one takes, should benefit from this work.

Formalization and Informalization

Formalization is the process of taking an informal piece of natural language mathematics, embedding it in into a theorem prover, constructing a model, and working with types instead of sets. This often requires significant amounts of work. We note some interesting artifacts about a piece of mathematics being formalized:

it may be formalized differently by two different people in many different ways

it may have to be modified, to include hidden lemmas, to correct of an error, or other bureaucratic obstacles

it may not type check, and only be presumed hypothetically to be 'a correct formalization' given evidence

Informalization, on the other hand is a process of taking a piece formal syntax, and turning it into a natural language utterance, along with commentary motivating

Category	Formal Proof	Informal Proof
Audience	Agda (and Human)	Human
Translation	Compiler	Human
Objectivity	Objective	Subjective
Historical	20th Century	\leq Euclid
Orientation	Syntax	Semantics
Inferability	Complete	Domain Expertise Necessary
Verification	PL Designer	Human
Ambiguity	Unambiguous	Ambiguous

Figure 6: Informal and Formal Proofs

and or relating it to other mathematics. It is a clarification of the meaning of a piece of code, suppressing certain details and sometimes redundantly reiterating other details. In figure Figure 8 we offer a few dimensions of comparison.

Mathematicians working in either direction know this is a respectable task, often leading to new methods, abstractions, and research altogether. And just as any type of machine translation, rule-based or statistical, on Virginia Woolf novel or Emily Dickinson poem from English to Mandarin would be absurd, so-to would the pretense that the methods we explore here using GF could actually match the competence of mathematicians translating work between a computer a book. Despite the futility of surpassing a mathematician at proof translation, it shouldn't deter those so inspired to try.

Syntactic Completeness and Semantic Adequacy

The GF pipeline, that of bidirectional translation through an intermediary abstract syntax representation, has two fundamental criteria that must be assessed for one to judge the success of an approach over both formalization and informalization.

The first criterion mentioned above, which we'll call *syntactic completeness*, says that a term either type-checks, or some natural language form can be deterministically transformed to a term that does type-check.

It asks the following : given an utterance or natural language expression that a mathematician might understand, does the GF grammar emit a well-formed syntactic expression in the target logic or programming language? The saying "grammars leak", can be transposed to say (NL) "proofs leak" in that they are certain to contain omissions.

This problem of syntactically complete mathematics is certain to be infeasible in many cases - a mathematician might not be able to reconstruct the unstated syntactic details of a proof in an discipline outside her expertise, it is at worthy pursuit to ask why it is so difficult! Additionally, certain inferable details may also detract from the natural language reading rather than assist. Perhaps most importantly, one does not know a priori that the generated expression in the logic has its intended meaning, other than through some meta verification procedure.

Conversely, given a well formed syntactic expression in, for instance, Agda, one can ask if the resulting English expression generated by GF is *semantically adequate*.

This notion of semantic adequacy is also delicate, as mathematicians themselves may dispute, for instance, the proof of a given proposition or the correct definition of some notion. However, if it is doubtful that there would be many mathematicians who would not understand some standard theorem statement and proof in an arbitrary introductory analysis text, even if one may dispute it's presentation, clarity, pedagogy, or other pedantic details. Whether one asks that semantic adequacy means some kind of sociological consensus among those with relevant expertise, or a more relaxed criterion that some expert herself understands the argument, a dubious perspective in scientific circles, semantic adequacy should appease at least one and potentially more mathematicians.




Figure 7: Formal and Informal Mathematics

We introduce these terms, syntactic completeness and semantic adequacy to highlight perspectives and insights that seems to underlie the biggest differences between informal and formal mathematics, as is show in Figure 9. We claim that mathematics, as done via a theorem prover, is a syntax oriented endeavor, whereas mathematics, as practiced by mathematicians, prioritizes semantic understanding. Developing a system which is able to formalize and informalize utterances which preserve syntactic completeness and semantic adequacy, respectively, is probably infeasible. Even introducing objective criteria to really judge these definitions is likely to be infeasible.

This perspective represents an observation and is not intended to judge whether the syntactic or semantic perspective on mathematics is better - there is a dialectical phenomena between the two. Let's highlight some advantages both provide, and try to distinguish more precisely what a syntactic and semantic perspective may be.

When the Agda user builds her proof, she is outsourcing much of the bookkeeping to the type-checker. This isn't purely a mechanical process though, she often does have to think, how her definitions will interact with downstream programs, as well as whether they are even sensible to begin with (i.e. does this have a proof). The syntactic side is expressly clear from the readers perspective as well. If Agda proofs were semantically coherent, one would only need to look at code, with perhaps a few occasional remarks about various intentions and conclusions, to understand the mathematics being expressed. Yet, papers are often written exclusively in Latex, where Agda proofs have to be reverse engineered, preserving only semantic details and forsaking syntactic nuance.

Oftentimes the code is kept in the appendix so as to provide a complete syntactic blueprint. But the act of writing an Agda proof and reading it is often orthogonal, as the term shadows the application of typing rules which enable its construction. The construction of the proof is entirely engaged with the types, whereas the human witness of a large term is either lost as to why it fulfills the typing judgment, she has to reexamine parts of the proof reasoning in her head or perhaps, try to rebuild interactively with Agda's help.

Even in cases where Agda code is included in a paper, it is most often the types which are emphasized and produced. Complex proof terms are seldom to be read on their own terms. The natural language description and commentary is still largely necessary to convey whatever results, regardless if the Agda code is self-contained. And while literate Agda is some type of bridge, it is still the commentary which in some sense unfolds the code and ultimately makes the Agda code legible.

This is particularly pronounced in the Coq programming language, where proof terms are built using Ltac, which can be seen as some kind of imperative syntactic metaprogramming over the core language, Gallina. The user rarely sees the internal proof tree that one becomes familiar with in Agda. The tactics are not typed, often feel very adhoc, and tacticals, sequences of tactics, may carry very little semantic value (or even possibly muddy one's understanding when reading proofs with unknown tactics). Indeed, since Ltac isn't itself typed, it often descends into the sorrows of so-called untyped languages (which are really uni-typed), and there are recent attempts to change this ? ?. From our perspective, the use of tactics is an additional syntactic obfuscation of what a proof should look like from the mathematicians perspective - and it is important to attempt to remedy this is. Alecyrton is one impressive development in giving Coq proofs more readability through a interactive back-end which shows the proof state, and offers other semantically appealing models like interactive graphics ?. This kind of system could and should inspire other proof assistants to allow for experimentation with syntactic alternative to linear code.

Tactics obviously have their uses, and sometimes enhance high level proof understanding, as tactics like *ring* or *omega* often save the reader overhead of parsing pedantic and uninformative details. For certain proofs, especially those involving many hundreds of cases, the metaprogramming facilities actually give one exclusive advantages not offered to the classical mathematician using pen and paper. Nonetheless, the dependent type theorist's dream that all mathematicians begin using theorem provers in their everyday work is largely just a dream, and with relatively little mainstream adoption by mathematicians, the future is all but clear.

Mathematicians may indeed like some of the facilities theorem provers provide, but ultimately, they may not see that as the "essence" of what they are doing. What is this essence? We will try to shine a small light on perhaps the most fundamental question in mathematics.

What is a proof?

A proof is what makes a judgment evident ?.

The proofs of Agda, and any programming language supporting proof development, are *formal proofs*. Formal proofs have no holes, and while there may very well be bugs in the underlying technologies supporting these proofs, formal proofs are seen as some kind of immutable form of data. One could say they provide *objective evidence* for judgments, which themselves are objective entities when encoded on a computer. What we call formal proofs might provide a science fiction writer an interesting thought experiment as regards communicating mathematics with an alien species incapable of understanding our language otherwise. Formal proofs, however, certainly don't appease all mathematicians writing for other mathematicians.

Mathematics, and the act of proving theorems, according to Brouwer is a social process. And because social processes between humans involve our linguistic faculties, we hope to elucidate what a proof with a simplified description. Suppose we have two humans, h_1 and h_2 . If h_1 claims to have a proof p_1 , and elaborates it to p_2 who claims she can either verify p_1 or reproduce and re-articulate it via p'_1 , such that h_1 and h_2 agree that p_1 and p'_1 are equivalent, then they have discovered some mathematics. In fact, in this guise mathematics, can be viewed as a science, even if in fact it is constructed instead of discovered.

An apt comparison is to see the mathematician as architect, whereas the computer scientist responsible for formalizing the mathematics is an engineer. The mathematics is the building which, like all human endeavors, is created via resources and labor of many people. The role of the architect is to envision the facade, the exterior layer directly perceived by others, giving a building its character, purpose, and function. The engineer is on the other hand, tasked with assuring the building gets built, doesn't collapse, and functions with many implicit features which the user of the building may not notice : the running water, insulation, and electricity. Whereas the architect is responsible for the building's *specification*, the engineer is tasked with its *implementation*.

We claim informal proofs are specifications and formal proofs are implementations. Additionally, via the propositions-as-types interpretation, one may see a logic as a specification and a PL as an implementation of a given logic, often with multiple ways of assigning terms to a given type. Therefore, one may see the mathematician ambiently developing a theorem in classical first order logic as providing a specification of a proposition in that language, whereas a given implementation of that theorem in Agda could be viewed as a model construction of some NL fragment, where truth in the model would correspond to termination of type-checking. Alternatively, during the informalization process, two different authors may suppress different details, or phrase a given utterance entirely differently, possibly leading to two different, but possibly similar proofs. Extrapolating our analogy, the same two architects given the same engineering plans could produce two entirely different looking and functioning buildings. Oftentimes though, it is the architect who has the vision, and the engineers who end up implementing the architect's art.

We also briefly explore the difference between the mathematician and the physicist. The physicist will often say under her breath to a class, “don’t tell anyone in the math department I’m doing this” when swapping an integral and a sum or other loose but effective tricks in her blackboard calculations. While there is an implicit assumption that there are theorems in analysis which may justify these calculations, it is not the physicist’s objective to be as rigorous as the mathematician. This is because the physicist is not using the mathematics as a syntactic mechanism to reflect the semantic domain of particles, energy, and other physical processes which the mathematics in physics serves to describe. The mathematician using Agda, needing to make syntactically complete arguments, needs to be obsessed with the details - whereas the “pen and paper” mathematician would need be reluctant to carry out all the excruciating syntactic details for sake of semantic clarity.

There isn’t a natural notion of equivalence between informal and formal proofs, but rather, loosely, some kind of adjunction between these two sets. We note the fact that the “acceptable” Natural language utterances aren’t inductively defined. This precludes us from actually constructing a canonical mathematical model of formal/informal relationship, but we most certainly believe that if the GF perspective of translation is used, there can at least be an approximation of what a model may look like. It is our contention that the linguist interested in the language of mathematics should perhaps be seen as a scientist, whose point is to contribute basic ideas and insights from which the architects and engineers can use to inform their designs.

Mathematicians seek model independence in their results (i.e., they don’t need a direct encoding of Fermat’s last theorem in set theory in order to trust its validity). This is one possible reason why there is so much reluctance to adopt proof assistant, because the implementation of a result in Coq, Agda, or HOL4 may lead to many permutations of the same result, each presumably representing the same piece of knowledge. It’s also noted a proof doesn’t obey the same universality that it does when it’s on paper or verbalized - that Agda 2.6.2, and its standard library, when updated in the future, may “break proofs”, as was seen in the introduction. While this is a unanimous problem with all software, we believe the GF approach offers at least a vision of not only linguistic, but also foundation agnosticism with respect to mathematics.

This thesis examines not just a practical problem, but touches many deep issues in some space in the intersection of the foundations, of mathematics, logic, computer science, and their relations studied via linguistic formalisms. These subjects, and their various relations, are the subject of countless hours of work and consideration by many great minds. We barely scratches the surface of a few of these developments, but it nonetheless, it is hoped, provides a nontrivial perspective at many important issues.

Recapitulating much of what was said, we hope that the following questions may have a new perspective :

- what are mathematical objects ?

- how do their encodings in different foundational formalisms affect their interpretations ?
- how does is mathematics develop as a social process ?
- how does what mathematics is and how it is done rely on given technologies of a given historical era. ?

While various branches of linguistics have seen rapid evolution due to, in large part, their adoption of mathematical tools, the dual application of linguistic tools to mathematics is quite sparse and open terrain. We hope the reader can walk away with an new appreciation to some of these questions and topics after reading this. These nuances we will not explore here, but shall be further elaborated in the future and and more importantly, hopefully inspire other readers to respond accordingly.

Although not given in specific detail, the view of what mathematics is, in both a philosophical and mathematical sense, as well as from the view of what a foundational perspective, requires deep consideration in its relation to linguistics. And while this work is perhaps just a finer grain of sandpaper on an incomplete and primordial marble sculpture, it is hoped that the sculptor's own reflection is a little bit more clear after we polish it here.

What is a proof revisited

Though philosophical discussion of visual thinking in mathematics has concentrated on its role in proof, visual thinking may be more valuable for discovery than proof ?

As an addendum to asking such a presumably simple question in the previous section, we briefly address the one particular oversimplification which was made. We briefly touch on what isn't just syntactic about mathematics, namely so-called "Proofs without Words" ? and other diagrammatic and visual reasoning tools generally. Because our work focuses on syntax, and is not generalized to other mathematical tools, we hope one considers this as well when pondering the language of mathematics.

The role of visualization in programming, logic, and mathematics generally offers an abundance of contrast to syntactically oriented alphanumeric alphabets, i.e. strings of symbols, which we discuss here. Although the trees in GF are visual, they are of intermediary form between strings in different languages, and therefore the type of syntax we're discussing here is strings, we hope a brief exploration of alternatives for concrete syntax will be fruitful. Targeting latex via GF for instance, is a small step in this direction.

Graphical Programming languages facilitating diagrammatic programming are one instance of a nonlinear syntax which would prove tricky but possible to implement via GF. Additionally, Globular, which allows one to carry out higher categorical constructions via globular sets is an interesting case study for a graphical programming language which is designed for theorem proving ?. Additionally, Ale-

cytron supports basic data structure visualization, like red-black trees which carry semantic content less easy in a string based-setting ?.

Visualization are ubiquitous in contemporary mathematics, whether it be analytic functions, knots, diagram chases in category theory, and a myriad of other visual tools which both assist understanding and inform our syntactic descriptions. We find these languages appealing because of their focus on a different kind of internal semantic sensation. The diagrammatic languages for monoidal categories, for example, also allow interpretations of formal proofs via topological deformations, and they have given semantic representations to various graphical languages like circuit diagrams and petri nets ?.

We also note that, while programming languages whose visual syntax evaluates to strings, means that all diagrams can in some sense be encoded in more traditional syntax, this is only for the computers sake - the human may consume the diagram as an abstract entity other than a string. There are often words to describe, but not to give visual intuition to many of the mathematical ideas we grasp. There are also, famously blind mathematicians who work in topology, geometry, and analysis ?. Bernard Morin, blinded at a young age, was a topologist who discovered the first eversion of a sphere by using clay models which were then diagrammatically transcribed by a colleague on the board. This remarkable use of something PL researchers would have a hard time imagining, in some sense, and warrants careful consideration of what the boundaries of proof assistants are capable of in terms of giving mathematicians more tangible constructions.

For if there is one message one should take away from this thesis, it is that there needs to be a coming to terms in the mathematics and TT communities, of the difference between *a proof*, both formal and informal, and the *the understanding of a proof*. The first is a mathematical judgment where one supplies evidence, via the form of a term that Agda can type-check and verify. A NL proof can be reviewed by a human. The understanding of a proof, however, is not done by anything but a human. And this internal understanding and processing of mathematical information, what I'll tongue-and-cheek call i-mathematics, with its externalization facilities being our main concerns in this thesis, requires much more work by future scholars.

Preliminaries

We give brief but relevant overviews of the background ideas and tools that went into the generation of this thesis.

Martin-Löf Type Theory

Judgments

With Kant, something important happened, namely, that the term judgement, Ger. Urteil, came to be used instead of proposition ?.

A central contribution of Per Martin-Löf in the development of type theory was the recognition of the centrality of judgments in logic. Many mathematicians aren't familiar with the spectrum of judgments available, and merely believe they are concerned with *the* notion of truth, namely *the truth* of a mathematical proposition or theorem. There are many judgments one can make which most mathematicians aren't aware of or at least never mention. Examples of both familiar and unfamiliar judgments include,

- A is true
- A is a proposition
- A is possible
- A is necessarily true
- A is true at time t

These judgments are understood not in the object language in which we state our propositions, possibilities, or probabilities, but as assertions in the metalanguage which require evidence for us to know and believe them. Most mathematicians may reach for their wallets if I come in and give a talk saying it is possible that the Riemann Hypothesis is true, partially because they already know that, and partially because it doesn't seem particularly interesting to say that something is possible, in the same way that a physicist may flinch if you say alchemy is possible. Most mathematicians, however, would agree that $P = NP$ is a proposition, and it is also possible, but isn't true.

For the logician these judgments may well be interesting because their may be logics in which the discussion of possibility or necessity is even more interesting than the discussion of truth. And for the type theorist interested in designing and building programming languages over many various logics, these judgments become a prime focus. The role of the type-checker in a programming language is to present evidence for, or decide the validity of the judgments. The four main judgments of type theory are given in natural language on the left and symbolically on the right.

- T is a type
- T and T' are equal types
- t is a term of type T
- t and t' are equal terms of type T
- $\vdash T$ type
- $\vdash T = T'$
- $\vdash t : T$
- $\vdash t = t' : T$

Frege's turnstile, \vdash , denotes a judgment.

These judgments become much more interesting when we add the ability for them to be interpreted in a some context with judgment hypotheses. Given a series of judgments J_1, \dots, J_n , denoted Γ , where J_i can depend on previously listed J' 's, we can make judgment J under the hypotheses, e.g. $J_1, \dots, J_n \vdash J$. Often these hypotheses J_i , alternatively called *antecedents*, denote variables which may occur freely in the *consequent* judgment J . For instance, the antecedent, $x : \mathbb{R}$ occurs freely in the syntactic expression $\sin x$, a which is given meaning in the judgment $\vdash \sin x : \mathbb{R}$. We write our hypothetical judgement as follows :

$$x : \mathbb{R} \vdash \sin x : \mathbb{R}$$

Rules

Martin-Löf systematically used the four fundamental judgments in the proof theoretic style of Prawitz. To this end, the intuitionistic formulation of the logical connectives just gives rules which admit an immediate computational interpretation. The main types of rules are type formation, introduction, elimination, and computation rules. The introduction rules for a type admit an induction principle derivable from that type's signature. Additionally, the β and η computation rules are derivable via the composition of introduction and elimination rules, which, if correctly formulated, should satisfy a relation known as harmony.

The fundamental notion of the lambda calculus, the function, is abstracted over a variable and returns a term of some type when applied to an argument which is subsequently reduced via the computational rules. Dependent Type Theory (DTT) generalizes this to allow the return type be parameterized by the variable being abstracted over. The dependent function forms the basis of the LF which underlies Agda and GF.

One reason why hypothetical judgments are so interesting is we can devise rules which allow us to translate from the metalanguage to the object language using lambda expressions. These play the role of a function in mathematics and implication in logic. This comes out in the following introduction rule :

Using this rule, we now see a typical judgment, typical in a field like from real analysis,

$$\vdash \lambda x. \sin x : \rightarrow$$

Equality :

Mathematicians denote this judgement

$$f: \mathbb{R} \rightarrow \mathbb{R}$$

$$x \mapsto \sin(x)$$

Propositions, Sets, and Types

While the rules of type theory have been well-articulated elsewhere, we provide briefly compare the syntax of mathematical constructions in FOL, one possible natural language use ?, and MLTT. From this vantage, these look like simple symbolic manipulations, and in some sense, one doesn't need a the expressive power of system like GF to parse these to the same form.

Additionally, it is worth comparing the type theoretic and natural language syntax with set theory. Now we bear witness to some deeper cracks than were visible above. We note that the type theoretic syntax is *the same* in both tables, whereas the set theoretic and logical syntax shares no overlap. This is because set theory and first order logic are treated as distinct domains, whereas in vanilla MLTT, there is no distinguishing mathematical types from logical types - everything is a type.

FOL	MLTT	NL FOL	NL MLTT
$\forall x P(x)$	$\Pi x : \tau. P(x)$	<i>for all x, p</i>	<i>the product over x in p</i>
$\exists x P(x)$	$\Sigma x : \tau. P(x)$	<i>there exists an x such that p</i>	<i>there exists an x in τ such that p</i>
$p \supset q$	$p \rightarrow q$	<i>if p then q</i>	<i>p to q</i>
$p \wedge q$	$p \times q$	<i>p and q</i>	<i>the product of p and q</i>
$p \vee q$	$p + q$	<i>p or q</i>	<i>the coproduct of p and q</i>
$\neg p$	$\neg p$	<i>it is not the case that p</i>	<i>not p</i>
\top	\top	<i>true</i>	<i>top</i>
\perp	\perp	<i>false</i>	<i>bottom</i>
$p = q$	$p \equiv q$	<i>p equals q</i>	<i>definitionally equal</i>

Figure 8: FOL vs MLTT

Set Theory	MLTT	NL Set Theory	NL MLTT
S	τ	<i>the set S</i>	<i>the type τ</i>
\mathbb{N}	Nat	<i>the set of natural numbers</i>	<i>the type nat</i>
$S \times T$	$S \times T$	<i>the product of S and T</i>	<i>the product of S and T</i>
$S \rightarrow T$	$S \rightarrow T$	<i>the function S to T</i>	<i>p to q</i>
$\{x P(x)\}$	$\Sigma x : \tau. P(x)$	<i>the set of x such that P</i>	<i>there exists an x in τ such that p</i>
\emptyset	\perp	<i>the empty set</i>	<i>bottom</i>
$?$	\top	<i>?</i>	<i>top</i>
$S \cup T$	$?$	<i>the union of S and T</i>	<i>?</i>
$S \subset T$	$S <: T$	<i>the subset S of T</i>	<i>S is a subtype of T</i>
$?$	U_1	<i>?</i>	<i>the second Universe</i>

Figure 9: Sets vs MLTT